

Harshita Dooja Poojary

harshita1.poojary@gmail.com | +1(213)-234-8695 | [in harshitapoojary](https://www.linkedin.com/in/harshitapoojary)

Solution-focused engineer with 4.5 years of experience in Python development, ML pipelines, and system integration, engaging stakeholders, translating complex concepts into actionable solutions, and optimizing client workflows through analytics and scalable platforms.

EDUCATION

University of Southern California, Los Angeles	Aug 2023 – May 2025
Master’s in Computer Science (Coursework- <i>Algorithms, Deep Learning, Artificial Intelligence, Machine Learning, Natural Language Processing</i>)	GPA: 3.81
University of Mumbai, India	July 2015 – May 2019
Bachelor’s in Computer Science Engineering (Coursework- <i>Software Engineering, Distributed Databases, Web Technologies</i>)	GPA: 3.83

SKILLS

Core Programming Languages: Python
Other Programming Languages: JavaScript, C++, Java, HTML/CSS, Scala
Frameworks & Libraries: Node.js, AngularJs, PySpark, PyTorch, Keras, TensorFlow, OpenCV, Streamlit
Databases: MongoDB, Redis, MySQL, Elasticsearch, Milvus, Hadoop, DynamoDB
Tools & Platforms: Docker, Kubernetes, AWS, Azure, Jenkins, GitHub, GitLab, JIRA, Postman, Swagger, PowerBI, Tableau, Google Analytics

WORK EXPERIENCE

AI Intern	July 2025 – Present
Deep Defense Solutions (Remote)	USA

- Developing ML pipelines in Python to detect false incident reports for the community management system and taking action on the incident.
- | | |
|---------------------|---------------------|
| Data Science Intern | Jun 2024 – Aug 2024 |
| Zero-True (Remote) | New York, USA |
- Collaborated with founders and client teams to design SQL/Python-integrated dashboards, presenting insights through interactive demos and advising on best practices for operational efficiency.
 - Implemented 5 Vuetify UI components that reduced client reporting errors by 15% as measured by accuracy metrics across accounts.
 - Optimized GPU inference efficiency by 30%, enabling scalable deployment of intensive workflows under constrained resources..

Student Researcher	Jan 2024 – Aug 2024
University of Southern California	Los Angeles, USA
<ul style="list-style-type: none">Under Prof. Seon Kim, analyzed object detection and classification models (YOLOv5/YOLOv8) on 10k Streetview imagery from Google Images and in-house dataset using data mining for tent detection and counting, optimizing performance in noisy urban conditions using geospatial metadata.Executed geospatial analysis at scale by coding grid-based segmentation logic to quantify encampment density across 502.7 square miles of Los Angeles, generating multi-intensity heatmaps using latitude-longitude clustering across five intensity levels.	

Machine Learning Engineer	Dec 2021 – Jul 2023
Reliance Jio	Mumbai, India
<ul style="list-style-type: none">Engineered and optimized RetinaFace detection models using TensorRT FP16 precision on T4 and A100 GPUs using C++ and Python, reducing face registration time by 40% and enabling real-time inference at scale across Jio’s parallel distributed clusters.Accelerated deployment velocity by 50% as measured by rollout metrics by automating CI/CD pipelines in Python with Docker, Kubernetes, GitLab, and JIRA for three ML systems—JioFace, Video Motion Detection, Number Plate Recognition.Delivered spoof detection with 90% accuracy as measured by production benchmarks by implementing ResNet/CDCNet models and scaling to 20+ sites.Directed 4 interns to curate datasets and validate models, ensuring reproducibility and deployment readiness across production environments.	

Software Engineer	Jun 2019 – Dec 2021
Reliance Jio	Mumbai, India
<ul style="list-style-type: none">As a part of the world's largest mobile data company, collaborated cross-functionally to build scalable microservices and REST APIs serving personalized ML-driven recommendations for 100M+ users, enhancing engagement by 20% using large-scale LLM models.Designed and optimized large-scale ingestion pipelines/templates for multi-device content delivery by crafting advanced SQL queries across fragmented relational datasets, improving throughput and aggregation performance.Analyzed user behaviors across 10 campaigns for marketing personalized advertisements, leveraging regression and time-series models to identify engagement patterns and optimize targeting strategies, which improved click-through rates and campaign ROI.Advised product managers and frontend teams on API integration and content delivery strategies, reducing latency by 15% as measured by system monitoring tools and boosting user engagement by 12% through streamlined backend-to-frontend data flows.	

Software Engineer	Jan 2019 – Jun 2019
TaksyKraft (Remote)	Hyderabad, India
<ul style="list-style-type: none">Developed engagement prediction APIs that improved ROI by 10% as measured by advertiser adoption by delivering real-time targeting solutions for 10+ clients.Engineered data pipelines and Tableau dashboards, reducing reporting delays by 40% as measured by turnaround times across enterprise accounts.	

PROJECTS

Unified Semantic Space for Multimodal Retrieval (PyTorch, Python, VectorDB)	
<ul style="list-style-type: none">Analyzed multimodal LLMs, including LLama-3, to answer questions over frames for MIT videos and using LangDB Vector Databases to store frame embeddings and retrieve user queries.Implemented a user interface using Streamlit to integrate ML models, enabling users to upload videos and query the system interactively.	
Analyzing NLP techniques on Low-Resource Languages(PyTorch)	
<ul style="list-style-type: none">Conducted idiom classification and metaphor classification in the low-resource Konkani language using BERT-based transformers, investigating performance-preserving attention head pruning.	
Comparison of Model Performance with Knowledge Distillation and Quantization (PyTorch, Python)	
<ul style="list-style-type: none">Fine-tuned LLama-3 model with 4-bit quantization using logit distillation with LoRa (Low-Rank Adaptation).Computed benchmark against GSM8k dataset with 8-shot learning to analyze the performance against the original model with 8B parameters.	